

# ArupChakraborty\_Assignment2.1

May 19, 2025

## 1 AWS Data Lake Query Notebook

This notebook sets up the environment, configures Athena, and executes various queries to analyze the ingested data from the AWS S3 Data Lake.

```
[2]: import boto3
import sagemaker
import pandas as pd

sess = sagemaker.Session()
bucket = sess.default_bucket()
role = sagemaker.get_execution_role()
region = boto3.Session().region_name
account_id = boto3.client("sts").get_caller_identity().get("Account")

sm = boto3.Session().client(service_name="sagemaker", region_name=region)
```

```
sagemaker.config INFO - Not applying SDK defaults from location:
/etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location:
/home/sagemaker-user/.config/sagemaker/config.yaml
```

```
[3]: s3_private_path_tsv = "s3://{}/amazon-reviews-pds/tsv".format(bucket)
print(s3_private_path_tsv)
```

```
s3://sagemaker-us-east-1-672518276407/amazon-reviews-pds/tsv
```

```
[4]: !aws s3 cp "dataset.csv" $s3_private_path_tsv/
```

```
upload: ./dataset.csv to s3://sagemaker-us-east-1-672518276407/amazon-reviews-
pds/tsv/dataset.csv
```

```
[5]: print(s3_private_path_tsv)
```

```
s3://sagemaker-us-east-1-672518276407/amazon-reviews-pds/tsv
```

```
[41]: import boto3
import pandas as pd
from io import StringIO
```

```

# S3 bucket and file path
bucket_name = "sagemaker-us-east-1-672518276407"
file_key = "amazon-reviews-pds/tsv/dataset.csv"

# Create an S3 client
s3 = boto3.client("s3")

try:
    # Fetch the CSV file from S3
    response = s3.get_object(Bucket=bucket_name, Key=file_key)

    # Read the content of the file
    content = response["Body"].read().decode("utf-8")

    # Load the content into a DataFrame
    data = pd.read_csv(StringIO(content))

    # Display the first few rows
    print(data.head())
except Exception as e:
    print(f"Error occurred: {e}")

```

Unnamed: 0	track_id	artists	\
0	0 5Su0ikwiRyPMVoIQDJUgSV	Gen Hoshino	
1	1 4qPNDBW1i3p13qLCtOKi3A	Ben Woodward	
2	2 1iJBSr7s7jYXzM8EGcbK5b	Ingrid Michaelson;ZAYN	
3	3 6lfxq3CG4xtTiEg7opyCyx	Kina Grannis	
4	4 5vjLSffimiIP26QG5WcN2K	Chord Overstreet	

  

	album_name	\
0	Comedy	
1	Ghost (Acoustic)	
2	To Begin Again	
3	Crazy Rich Asians (Original Motion Picture Sou...	
4	Hold On	

  

	track_name	popularity	duration_ms	explicit	\
0	Comedy	73	230666	False	
1	Ghost - Acoustic	55	149610	False	
2	To Begin Again	57	210826	False	
3	Can't Help Falling In Love	71	201933	False	
4	Hold On	82	198853	False	

  

	danceability	energy	...	loudness	mode	speechiness	acousticness	\
0	0.676	0.4610	...	-6.746	0	0.1430	0.0322	
1	0.420	0.1660	...	-17.235	1	0.0763	0.9240	

2	0.438	0.3590	...	-9.734	1	0.0557	0.2100
3	0.266	0.0596	...	-18.515	1	0.0363	0.9050
4	0.618	0.4430	...	-9.681	1	0.0526	0.4690

  

	instrumentalness	liveness	valence	tempo	time_signature	track_genre
0	0.000001	0.3580	0.715	87.917	4	acoustic
1	0.000006	0.1010	0.267	77.489	4	acoustic
2	0.000000	0.1170	0.120	76.332	4	acoustic
3	0.000071	0.1320	0.143	181.740	3	acoustic
4	0.000000	0.0829	0.167	119.949	4	acoustic

[5 rows x 21 columns]

```
[43]: # Filter songs with popularity greater than or equal to 99
df_high_popularity = data[data['popularity'] >= 99]
df_high_popularity = df_high_popularity[['artists', 'track_name', 'popularity']]
print("Songs with popularity 99:")
print(df_high_popularity)
```

Songs with popularity 99:

	artists	track_name	popularity
20001	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100
51664	Bizarrap;Quevedo	Quevedo: Bzrp Music Sessions, Vol. 52	99
81051	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100

```
[44]: # Group by artist and calculate the average popularity
df_avg_popularity = data.groupby('artists')['popularity'].mean().reset_index()

# Filter artists with an average popularity of 92
df_avg_popularity_92 = df_avg_popularity[df_avg_popularity['popularity'] == 92]
print("Artists with an average popularity of 92:")
print(df_avg_popularity_92)
```

Artists with an average popularity of 92:

	artists	popularity
11491	Harry Styles	92.0
22845	Rema;Selena Gomez	92.0

```
[46]: # Group by genre and calculate the average energy
df_avg_energy = data.groupby('track_genre')['energy'].mean().reset_index()

# Get the top 10 genres by average energy
df_top10_genres = df_avg_energy.sort_values(by='energy', ascending=False).
    .head(10)
print("Top 10 genres with the highest average energy:")
print(df_top10_genres)
```

Top 10 genres with the highest average energy:

	track_genre	energy
22	death-metal	0.931470
42	grindcore	0.924201
72	metalcore	0.914485
46	happy	0.910971
49	hardstyle	0.901246
27	drum-and-bass	0.876635
6	black-metal	0.874897
50	heavy-metal	0.874003
78	party	0.871237
61	j-idol	0.868677

```
[47]: # Filter tracks by the artist 'Bad Bunny'
bad_bunny_count = data[data['artists'] == 'Bad Bunny'].shape[0]
print(f"Number of tracks featuring Bad Bunny: {bad_bunny_count}")
```

Number of tracks featuring Bad Bunny: 48

```
[48]: # Group by genre and find the maximum popularity within each genre
df_genre_popularity = data.groupby('track_genre')['popularity'].max().
    ↪reset_index()

# Sort by popularity and get the top 10 genres
df_top10_popular_genres = df_genre_popularity.sort_values(by='popularity',
    ↪ascending=False).head(10)
print("Top 10 genres by their most popular track:")
print(df_top10_popular_genres)
```

Top 10 genres by their most popular track:

	track_genre	popularity
80	pop	100
20	dance	100
51	hip-hop	99
68	latino	98
89	reggaeton	98
30	edm	98
67	latin	98
88	reggae	98
79	piano	96
90	rock	96

```
[39]: statement = "SHOW TABLES in {}".format('dsoaws')

df_show = pd.read_sql(statement, conn)
df_show.head(5)
```

/tmp/ipykernel\_3285/2205842015.py:3: UserWarning: pandas only supports SQLAlchemy connectable (engine/connection) or database string URI or sqlite3

DBAPI2 connection. Other DBAPI2 objects are not tested. Please consider using SQLAlchemy.

```
df_show = pd.read_sql(statement, conn)
```

```
[39]:          tab_name
0  amazon_reviews_tsv
```

```
[34]: # SageMaker session setup
sess = sagemaker.Session()
bucket = sess.default_bucket()
role = sagemaker.get_execution_role()
region = boto3.Session().region_name

# Athena connection setup
database_name = "dsoaws"
table_name_tsv = "amazon_reviews_tsv"

conn = connect(region_name=region, s3_staging_dir=s3_staging_dir)
print("Setup completed.")
```

Setup completed.

```
[61]: # Athena connection setup
conn = connect(region_name=region, s3_staging_dir=s3_staging_dir)

# Improved execute_query function
def execute_query(query, description="Query Result"):
    """
    Executes a Pandas query or calculation and prints the result.

    Parameters:
        query (pd.DataFrame or int): The result of the query.
        description (str): A brief description of the query result.
    """
    try:
        print(f"\n{description}")
        print("=" * len(description))

        # Check if the query returned a DataFrame
        if isinstance(query, pd.DataFrame):
            if not query.empty:
                display(query.head(100)) # Display the first 100 rows
                print(f"Total Records: {len(query)}")
            else:
                print("No matching records found.")
        # Check if the query returned a single value (like a count)
        elif isinstance(query, int):
            print(f"Result: {query}")
```

```

    else:
        print("Unexpected query result type.")
except Exception as e:
    print(f"Error executing query: {e}")

```

### 1.1 List artist, track\_name, and popularity for songs with popularity 99

```

[62]: import pandasql as psql

# SQL-like query on the DataFrame
query = """
SELECT artists, track_name, popularity
FROM data
WHERE popularity >= 99
"""
result = psql.sqldf(query, locals())
execute_query(result, "Songs with Popularity 99 (SQL-like)")

```

Songs with Popularity 99 (SQL-like)

=====

	artists	track_name	popularity
0	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100
1	Bizarrap;Quevedo	Quevedo: Bzrp Music Sessions, Vol. 52	99
2	Sam Smith;Kim Petras	Unholy (feat. Kim Petras)	100

Total Records: 3

### 1.2 List artists with an average popularity of 92

```

[65]: query = """
SELECT artists, AVG(popularity) AS avg_popularity
FROM data
GROUP BY artists
HAVING AVG(popularity) = 92
"""
result = psql.sqldf(query, locals())
execute_query(result, "List artists with an average popularity of 92")

```

List artists with an average popularity of 92

=====

	artists	avg_popularity
0	Harry Styles	92.0
1	Rema;Selena Gomez	92.0

Total Records: 2

### 1.3 List the Top 10 genres with the highest average energy

```
[66]: query = """
SELECT track_genre, AVG(energy) AS avg_energy
FROM data
GROUP BY track_genre
ORDER BY avg_energy DESC
LIMIT 10
"""
result = psql.sqldf(query, locals())
execute_query(result, "List the Top 10 genres with the highest average energy")
```

List the Top 10 genres with the highest average energy

=====

	track_genre	avg_energy
0	death-metal	0.931470
1	grindcore	0.924201
2	metalcore	0.914485
3	happy	0.910971
4	hardstyle	0.901246
5	drum-and-bass	0.876635
6	black-metal	0.874897
7	heavy-metal	0.874003
8	party	0.871237
9	j-idol	0.868677

Total Records: 10

### 1.4 How many tracks is Bad Bunny on?

```
[67]: query = """
SELECT COUNT(*)
FROM data
WHERE artists = 'Bad Bunny'
"""
result = psql.sqldf(query, locals())
execute_query(result, "How many tracks is Bad Bunny on?")
```

How many tracks is Bad Bunny on?

=====

	COUNT(*)
0	48

Total Records: 1

### 1.5 Show the top 10 genres in terms of popularity, sorted by their most popular track

```
[68]: query = """
SELECT track_genre, MAX(popularity) AS max_popularity
FROM data
GROUP BY track_genre
ORDER BY max_popularity DESC
LIMIT 10
"""
result = psql.sqldf(query, locals())
execute_query(result, "Show the top 10 genres in terms of popularity, sorted by_
↳their most popular track")
```

Show the top 10 genres in terms of popularity, sorted by their most popular track

=====

=

	track_genre	max_popularity
0	pop	100
1	dance	100
2	hip-hop	99
3	reggaeton	98
4	reggae	98
5	latino	98
6	latin	98
7	edm	98
8	rock	96
9	piano	96

Total Records: 10

[ ]: